

Managing HPC Software Complexity with Spack

The most recent version of these slides can be found at:
<https://spack.readthedocs.io/en/latest/tutorial.html>

Supercomputing 2018 Full-day Tutorial
November 12, 2018
Dallas, Texas



LLNL-PRES-806064

This work was performed under the auspices of the U.S. Department of Energy by Lawrence Livermore National Laboratory under contract DE-AC52-07NA27344.
Lawrence Livermore National Security, LLC.

github.com/spack/spack

 Lawrence Livermore
National Laboratory

 SCITAS
Software Configuration and Installation Tools

 EPFL
ÉCOLE POLYTECHNIQUE
FÉDÉRALE DE LAUSANNE

 NERSC

 I

 SC18
Dallas, TX

Tutorial Materials

Materials: Download the latest version of slides and handouts at:

<http://spack.readthedocs.io>

Slides and hands-on scripts are in the “Tutorial” Section of the docs.

- Spack GitHub repository: <http://github.com/spack/spack>
- Spack Documentation: <http://spack.readthedocs.io>



Tweet at us!

@spackpm #SC18

Tutorial Presenters



**Lawrence Livermore
National Laboratory**

Todd Gamblin
Gregory Becker
Greg Lee
Matt Legendre



Mario Melara



University of Illinois

Adam Stewart

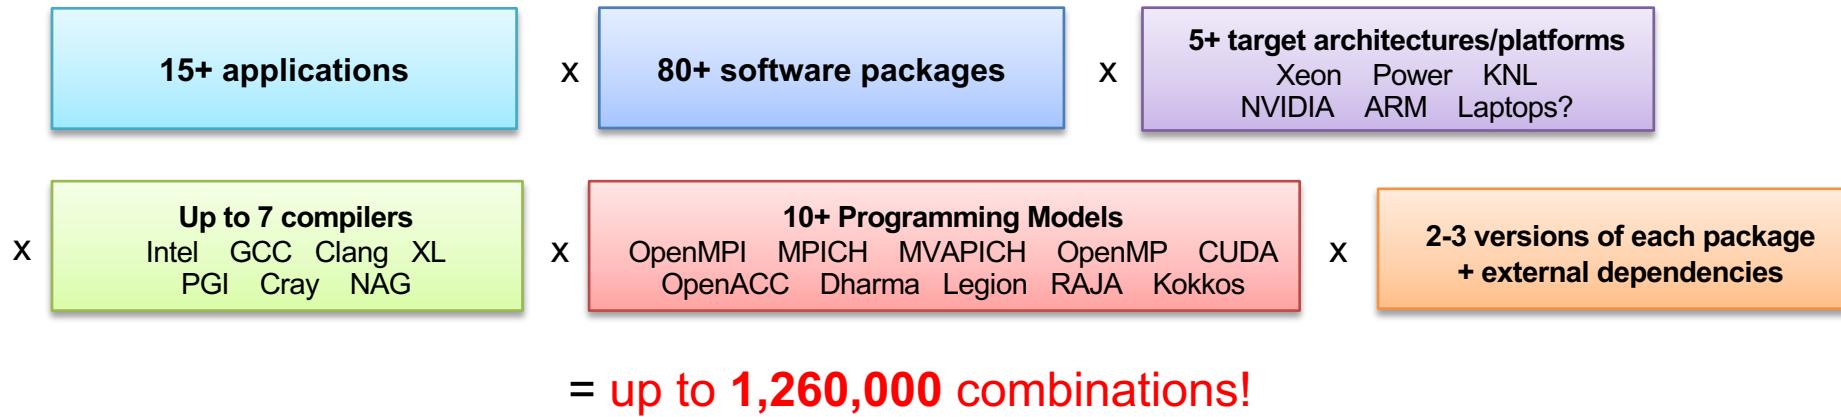


ÉCOLE POLYTECHNIQUE
FÉDÉRALE DE LAUSANNE



Massimiliano Culpo

The complexity of the exascale ecosystem threatens productivity.



- Every application has its own stack of dependencies.
- Developers, users, and facilities dedicate (many) FTEs to building & porting.
- Often trade reuse and usability for performance.

We must make it easier to rely on others' software!

What is the “production” environment for HPC?

- Someone’s home directory?
- LLNL? LANL? Sandia? ANL? LBL? TACC?
 - Environments at large-scale sites are very different.
- Which MPI implementation?
- Which compiler?
- Which dependencies?
- Which versions of dependencies?
 - Many applications require specific dependency versions.



Real answer: there isn’t a single production environment or a standard way to build.
Reusing someone else’s software is HARD.

Spack is a flexible package manager for HPC

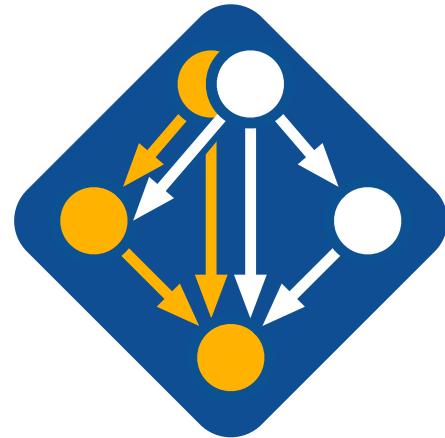
- How to install Spack:

```
$ git clone https://github.com/spack/spack
$ . spack/share/spack/setup-env.sh
```

- How to install a package:

```
$ spack install hdf5
```

- HDF5 and its dependencies are installed within the Spack directory.
- Unlike typical package managers, Spack can also install many variants of the same build.
 - Different compilers
 - Different MPI implementations
 - Different build options



Get Spack!

<http://github.com/spack/spack>



@spackpm

Who can use Spack?

People who want to use or distribute software for HPC!

1. End Users of HPC Software

- Install and run HPC applications and tools

2. HPC Application Teams

- Manage third-party dependency libraries

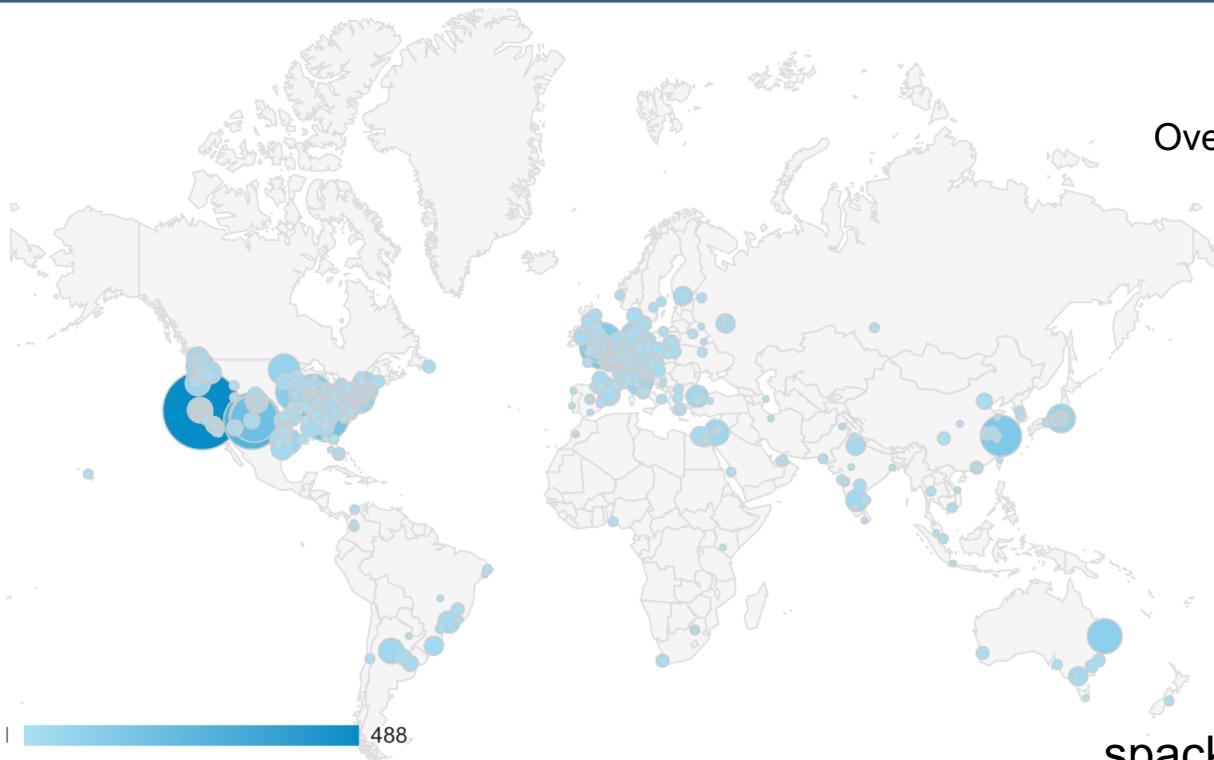
3. Package Developers

- People who want to package their own software for distribution

4. User support teams at HPC Centers

- People who deploy software for users at large HPC sites

Spack is used worldwide!



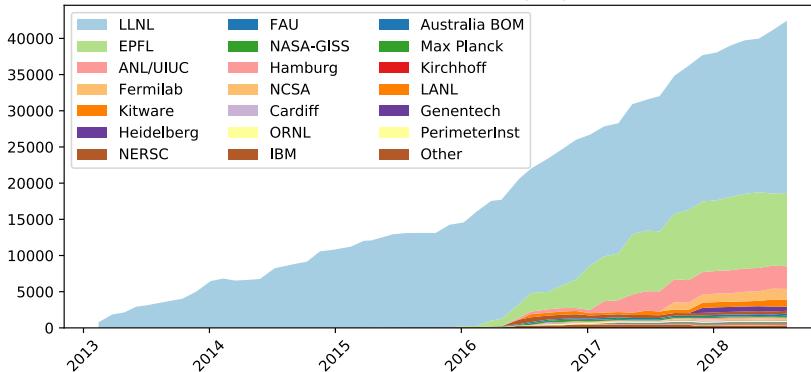
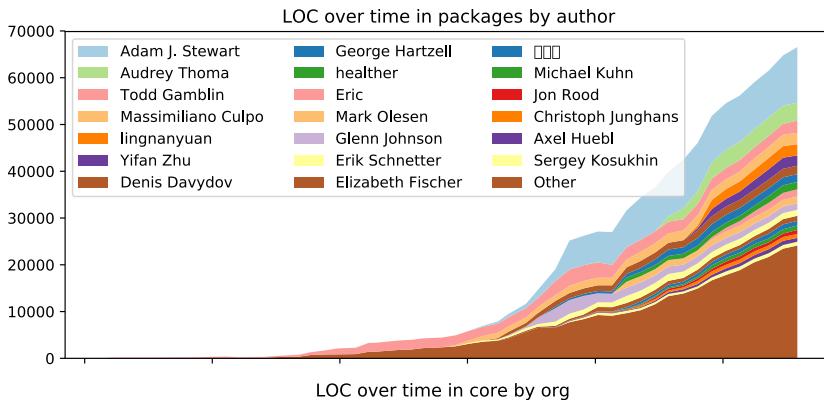
Over **2,900** software packages

Over **150,000** downloads in the past year

Over **300** contributors
from labs, academia, industry

Plot shows sessions on
`spack.readthedocs.io` for one month

Contributions to Spack continue to grow!



- In November 2015, LLNL provided most of the contributions to Spack
- Since then, we've gone from 300 to over 2,900 packages
- Most packages are from external contributors!
- Many contributions in core, as well.
- We are committed to sustaining Spack's open source ecosystem!



The Spack community now spans DOE and beyond

- **30+ organizations
140+ contributors**

Sharing over 2,900 packages and growing

- **Other use cases:**

- ARM using for entire compiler regression suite.
- LIGO collaboration using for deployment
- Intel using Spack to package ML software
- NERSC using Spack on Cori: Cray support.
- EPFL (Switzerland) contributing core features.
- Fermi, CERN, BNL: high energy physics.
- ANL using Spack on production Linux clusters.
- NASA packaging an Ice model code.
- ORNL working with us on Spack for CORAL.
- Kitware: core features, ParaView, Qt, UV-CDAT support



Spack v0.12.0 was just released

- **New stuff:**
 1. Spack environments (covered today)
 2. `spack.yaml` and `spack.lock` files for tracking dependencies (covered today)
 3. Custom configurations via command line (covered today)
 4. Better support for linking Python packages into view directories (pip in views)
 5. Packages have more control over compiler flags via flag handlers
 6. Better support for module file generation
 7. Better support for Intel compilers, Intel MPI, etc.
 8. Many performance improvements, improved startup time
- Spack is now permissively licensed under **Apache-2.0 or MIT**
 - previously LGPL
- Over 2,900 packages (800 added since last year)

Related Work

Spack is not the first tool to automate builds

- Inspired by copious prior work

1. “Functional” Package Managers

- Nix
- GNU Guix

<https://nixos.org/>

<https://www.gnu.org/s/guix/>

2. Build-from-source Package Managers

- Homebrew
- MacPorts

<http://brew.sh>

<https://www.macports.org>

Other tools in the HPC Space:

▪ Easybuild

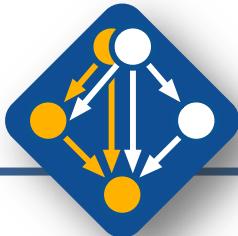
- An *installation* tool for HPC
- Focused on HPC system administrators – different package model from Spack
- Relies on a fixed software stack – harder to tweak recipes for experimentation

<http://hpcugent.github.io/easybuild/>

▪ Conda

- Very popular binary package manager for data science
- Not targeted at HPC; generally unoptimized binaries

<https://conda.io>



Other Spack events at SC18

Today, C148: This tutorial!

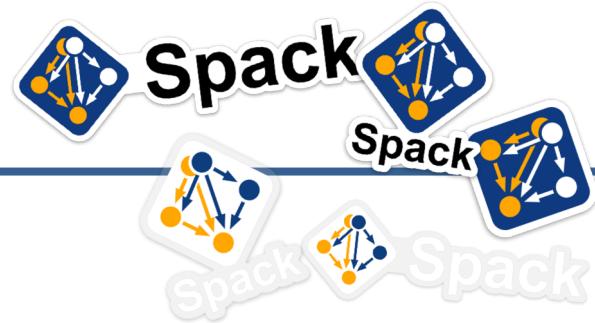
Tuesday 12:15pm – 1:15pm, D227

- Spack Community BOF

Wednesday 3:00pm – 4:00pm at the DOE Booth

- Spack Roundtable: Come meet the Spack Developers!
- Discuss recent features, plans, and issues.

Join us at these events for more!



We have shirts and stickers!

Tutorial Overview (times are estimates)

1.	For everyone:	Welcome & Overview	8:30 - 8:40
2.	For everyone:	Basics of Building and Linking	8:40 - 9:00
3.	For everyone:	Spack Basics (hands on)	9:00 - 9:45
4.	For everyone:	Core Spack Concepts	9:45 - 10:00
		-- Break --	10:00 - 10:30
5.	For everyone:	Configuration (hands-on)	10:30 - 11:00
6.	For packagers:	Making your own Packages (hands on)	11:00 - 12:00
		-- Tutorials Lunch (C1/2/3/4 Ballroom) --	12:00 - 1:30
7.	For everyone:	Spack environments, spack.yaml/spack.lock (hands-on)	1:30 - 2:15
8.	For HPC centers:	Generating Module Files with Spack (hands-on)	2:15 - 3:00
		-- Coffee Break --	3:30 - 3:30
9.	For packagers:	Advanced Packaging (build systems)	3:30 - 3:50
10.	For packagers:	Advanced Packaging (virtual dependencies: BLAS/LAPACK/MPI)	3:50 - 4:20
11.	For everyone:	Build your own packages with help from the Spack team!	4:20 - 5:00

Building & Linking Basics

What's a package manager?

- Spack is a ***package manager***
 - Does not replace Cmake/Autotools
 - Packages built by Spack can have any build system they want
 - Spack manages ***dependencies***
 - Drives package-level build systems
 - Ensures consistent builds
 - Determining magic configure lines takes time
 - Spack is a cache of recipes
- Package Manager
- High Level Build System
- Low Level Build System
- Manages package installation
 - Manages dependency relationships
 - Drives package-level build systems
- Cmake, Autotools
 - Handle library abstractions
 - Generate Makefiles, etc.
- Make, Ninja
 - Handles dependencies among commands in a single build

Static vs. shared libraries

- Static libraries: `libfoo.a`
 - .a files are archives of .o files (object files)
 - Linker includes needed parts of a static library in the output executable
 - No need to find dependencies at runtime – only at build time.
 - Can lead to large executables
 - Often hard to build a completely static executable on modern systems.
- Shared libraries: `libfoo.so` (Linux), `libfoo.dylib` (MacOS)
 - More complex build semantics, typically handled by the build system
 - Must be found by `ld.so` or `dyld` (dynamic linker) and loaded at runtime
 - Can cause lots of headaches with multiple versions
 - 2 main ways:
 - `LD_LIBRARY_PATH`: environment variable configured by user and/or module system
 - `RPATH`: paths embedded in executables and libraries, so that they know where to find their own dependencies.

API and ABI Compatibility

▪ API: Application Programming Interface

- Source code functions and symbol names exposed by a library
- If API of a dependency is backward compatible, source code need not be changed to use it
- **May** need to recompile code to use a new version.

▪ ABI: Application Binary Interface

- Calling conventions, register semantics, exception handling, etc.
- Defined by how the compiler builds a library
 - Binaries generated by different compilers are typically ABI-incompatible.
- May also include things like standard runtime libraries and compiler intrinsic functions
- May also include values of hard-coded symbols/constants in headers.

▪ HPC code, including MPI, is typically API-compatible but not ABI-compatible.

- Causes many build problems, especially for dynamically loaded libraries
- Often need to rebuild to get around ABI problems
- Leads to combinatorial builds of software at HPC sites.

3 major build systems to be aware of

1. Make (usually GNU Make)

- <https://www.gnu.org/software/make/>

2. GNU Autotools

- Automake: <https://www.gnu.org/software/automake/>
- Autoconf: <https://www.gnu.org/software/autoconf/>
- Libtool: <https://www.gnu.org/software/libtool/>

3. CMake:

- <https://cmake.org>

Spack has built-in support for these plus Waf, Perl, Python, and R

Make and GNU Make

- Many projects opt to write their own Makefiles.
 - Can range from simple to very complicated
- Make declares some standard variables for various compilers
 - Many HPC projects don't respect them
 - No standard install prefix convention
 - Makefiles may not have install target
- Automating builds with Make usually requires editing files
 - Typical to use sed/awk/some other regular expression tool on Makefile
 - Can also use patches

github.com/spack/spack

Typical build incantation

```
<edit Makefile>  
make PREFIX=/path/to/prefix
```

Configure options

None. Typically must edit Makefiles.

Environment variables

CC	CFLAGS	LDFLAGS
CXX	CXXFLAGS	LIBS
FC	FFLAGS	CPP
F77		

Autotools

- Three parts of autotools:
 - autoconf: generates a portable configure script to inspect build host
 - automake: high-level syntax for generating lower-level Makefiles.
 - libtool: abstraction for shared libraries
- Typical variables are similar to make
- Much more consistency among autotools projects
 - Wide use of standard variables and configure options
 - Standard install target, staging conventions.

Typical build incantation

```
./configure --prefix=/path/to/install_dir  
make  
make install
```

Configure options

```
./configure \  
--prefix=/path/to/install_dir \  
--with-package=/path/to/dependency \  
--enable-foo \  
--disable-bar
```

Environment variables

CC	CFLAGS	LDFLAGS
CXX	CXXFLAGS	LIBS
FC	FFLAGS	CPP
F77		

CMake

- Gaining popularity
- Arguably easier to use (for developers) than autotools
- Similar standard options to autotools
 - different variable names
 - More configuration options
 - Abstracts platform-specific details of shared libraries
- Most CMake projects should be built “out of source”
 - Separate build directory from source directory

Typical build incantation

```
mkdir BUILD && cd BUILD  
cmake -DCMAKE_INSTALL_PREFIX=/path/to/install_dir ..  
make  
make install
```

Configure options

```
cmake \  
  -D CMAKE_INSTALL_PREFIX=/path/to/install_dir \  
  -D ENABLE_FOO=yes \  
  -D ENABLE_BAR=no \  
  ..
```

Common –D options

CMAKE_C_COMPILER	CMAKE_C_FLAGS
CMAKE_CXX_COMPILER	CMAKE_CXX_FLAGS
CMAKE_Fortran_COMPILER	CMAKE_Fortran_FLAGS
CMAKE_SHARED_LINKER_FLAGS	CMAKE_EXE_LINKER_FLAGS
CMAKE_STATIC_LINKER_FLAGS	

Spack Basics

Spack provides a *spec* syntax to describe customized DAG configurations

```
$ spack install mpileaks           unconstrained
$ spack install mpileaks@3.3       @ custom version
$ spack install mpileaks@3.3 %gcc@4.7.3    % custom compiler
$ spack install mpileaks@3.3 %gcc@4.7.3 +threads  +/- build option
$ spack install mpileaks@3.3 cppflags="-O3 -g3"   setting compiler flags
$ spack install mpileaks@3.3 os=CNL10 target=haswell setting target for X-compile
$ spack install mpileaks@3.3 ^mpich@3.2 %gcc@4.9.3  ^ dependency information
```

- Each expression is a *spec* for a particular configuration
 - Each clause adds a constraint to the spec
 - Constraints are optional – specify only what you need.
 - Customize install on the command line!
- Spec syntax is recursive
 - Full control over the combinatorial build space

`spack list` shows what packages are available

```
$ spack list
==> 303 packages.
activeharmony cgal fish gtkplus libgd mesa openmpi py-coverage py-pycparser qt tcl
adept-utils cgm flex harfbuzz libggp-error metis openspeedshop py-cython py-pyelftools qthreads texinfo
apex cityhash fltk hdf libjpeg-turbo Mitos openssl py-dateutil py-pygments R the_silver_searcher
arpack cleverleaf flux hdfs libjson-c mpc otf py-epydoc py-pylint ravel thrift
asciidoc cloog fontconfig hwloc libmng mpe2 otf2 py-funcsigs py-pypar readline tk
atk cmake freetype hybre libmonitor mpfr pango py-genders py-pyparsing rose tmux
atlas cmocka gasnet icu libNBc mpibash papi py-gnuplot py-pyqt rsync tmuxinator
atop coreutils gcc icu4c libpicaaccess mpich paraver py-hspy py-pyside ruby trilinos
autoconf cppcheck gdb ImageMagick libpng mpileaks paraview py-ipython py-pytables SAMRAI uncrustify
automated cram gdk-pixbuf isl libsodium mrnet parmetis py-libxml2 py-python-daemon samtools util-linux
automake cscope geos jdk libtiff mumps parpack py-lockfile py-ptz scalasca valgrind
bear cube gflags jemalloc libtool munge patchelf py-mako py-rpy2 scorep vim
bib2xhtml curl ghostscript jpeg libunwind muster pcre py-matplotlib py-scientificpython scotch vtk
binutils czmq git judy libuuid mvapich2 pcre2 py-mock py-scikit-learn scr wget
bison damselfly glib julia libxcb nasm pdt py-mpi4py py-scipy silo wx
boost dbus glm launchmon libxml2 ncdu petsc py-mx py-setuptools snappy wxpropgrid
bowtie2 docbook-xml global lcms libxshmfence ncurses pidx py-mysqldb1 py-shiboken sparsehash xcb-proto
boxlib doxygen glog leveldb libxslt netcdf pixman py-nose py-sip spindle xerces-c
bzip2 driproto glpk libarchive llvm netgauge pkg-config py-numexpr py-six spot xz
cairo dtcmpl gmp libcerf llvm-lld netlib-blas pmgr_collective py-numpy py-sphinx sqlite yasm
callpath dyninst gmsk libcircle lmdb netlib-lapack postgresql py-pandas py-sympy stat zeromq
cblas eigen gnuplot libdrm lmod netlib-scalapack ppl py-pbr py-periodictable py-twisted sundials zlib
cbtf elfutils gnutls libdwarf lua nettle protobuf py-pexpect py-urwid szip tar
cbtf-argonavis elpa gperf libedit lwgrp ninja py-astropy py-baseemap py-pil py-virtualenv taskd tau
cbtf-krell expat gperf tools libelf lwm2 ompss py-basemap py-biopython py-pillow py-yapf
cbtf-lanl extrae graphlib libevent matio ompt-openmp py-blessings py-pmw python
cereal exuberant-ctags graphviz libffl mbedTLS opari2 py-cffi py-pychecker qhull
cfitsio fftw gsl libgcrypt memaxes openblas
```

- Spack has over 2,900 packages now.

`spack find` shows what is installed

```
$ spack find
==> 103 installed packages.
-- linux-rhel6-x86_64 / gcc@4.4.7 -----
ImageMagick@6.8.9-10 glib@2.42.1 libtiff@4.0.3 pango@1.36.8 qt@4.8.6
SAMRAI@3.9.1 graphlib@2.0.0 libtool@2.4.2 parmetis@4.0.3 qt@5.4.0
adept-utils@1.0 gtkplus@2.24.25 libxcb@1.11 pixman@0.32.6 ravel@1.0.0
atk@2.14.0 harfbuzz@0.9.37 libxml2@2.9.2 py-dateutil@2.4.0 readline@6.3
boost@1.55.0 hdf5@1.8.13 llvm@3.0 py-ipython@2.3.1 scotch@6.0.3
cairo@1.14.0 icu@54.1 metis@5.1.0 py-nose@1.3.4 starpu@1.1.4
callpath@1.0.2 jpeg@9a mpich@3.0.4 py-numumpy@1.9.1 stat@2.1.0
dyninst@8.1.2 libdwarf@20130729 ncurses@5.9 py-pytz@2014.10 xz@5.2.0
dyninst@8.1.2 libelf@0.8.13 ocr@2015-02-16 py-setup-tools@11.3.1 zlib@1.2.8
fontconfig@2.11.1 libffi@3.1 openssl@1.0.1h py-six@1.9.0
freetype@2.5.3 libmng@2.0.2 otf@1.12.5salmon python@2.7.8
gdk-pixbuf@2.31.2 libpng@1.6.16 otf2@1.4 qhull@1.0

-- linux-rhel6-x86_64 / gcc@4.8.2 -----
adept-utils@1.0.1 boost@1.55.0 cmake@5.6-special libdwarf@20130729 mpich@3.0.4
adept-utils@1.0.1 cmake@5.6 dyninst@8.1.2 libelf@0.8.13 openmpi@1.8.2

-- linux-rhel6-x86_64 / intel@14.0.2 -----
hwloc@1.9 mpich@3.0.4 starpu@1.1.4

-- linux-rhel6-x86_64 / intel@15.0.0 -----
adept-utils@1.0.1 boost@1.55.0 libdwarf@20130729 libelf@0.8.13 mpich@3.0.4

-- linux-rhel6-x86_64 / intel@15.0.1 -----
adept-utils@1.0.1 callpath@1.0.2 libdwarf@20130729 mpich@3.0.4
boost@1.55.0 hwloc@1.9 libelf@0.8.13 starpu@1.1.4
```

- All the versions coexist!
 - Multiple versions of same package are ok.
- Packages are installed to automatically find correct dependencies.
- Binaries work *regardless of user's environment*.
- Spack also generates module files.
 - Don't have to use them.

Users can query the full dependency configuration of installed packages.

```
$ spack find callpath  
==> 2 installed packages.  
-- linux-rhel6-x86_64 / clang@3.4 -----  
callpath@1.0.2  
-- linux-rhel6-x86_64 / gcc@4.9.2 -----  
callpath@1.0.2
```



Expand dependencies with `spack find -d`

```
$ spack find -dl callpath  
==> 2 installed packages.  
-- linux-rhel6-x86_64 / clang@3.4 -----  
xv2clz2    callpath@1.0.2  
ckjazss    ^adept-utils@1.0.1  
3ws43m4     ^boost@1.59.0  
ft7znm6    ^mpich@3.1.4  
qqnuet3    ^dyninst@8.2.1  
3ws43m4     ^boost@1.59.0  
g65rdud    ^libdwarf@20130729  
cj5p5fk    ^libelf@0.8.13  
cj5p5fk    ^libelf@0.8.13  
g65rdud    ^libdwarf@20130729  
cj5p5fk    ^libelf@0.8.13  
cj5p5fk    ^libelf@0.8.13  
ft7znm6    ^mpich@3.1.4  
-- linux-rhel6-x86_64 / gcc@4.9.2 -----  
udltshs    callpath@1.0.2  
rfsu7fb    ^adept-utils@1.0.1  
ybet64y    ^boost@1.55.0  
aa4ar6i    ^mpich@3.1.4  
tmmnge5    ^dyninst@8.2.1  
ybet64y    ^boost@1.55.0  
g2mxrl2    ^libdwarf@20130729  
ynpai3j    ^libelf@0.8.13  
ynpai3j    ^libelf@0.8.13  
g2mxrl2    ^libdwarf@20130729  
ynpai3j    ^libelf@0.8.13  
ynpai3j    ^libelf@0.8.13  
aa4ar6i    ^mpich@3.1.4
```

- Architecture, compiler, versions, and variants may differ between builds.

Spack manages installed compilers

- Compilers are automatically detected
 - Automatic detection determined by OS
 - Linux: PATH
 - Cray: `module avail`
- Compilers can be manually added
 - Including Spack-built compilers

```
$ spack compilers
==> Available compilers
-- gcc -----
gcc@4.2.1      gcc@4.9.3

-- clang -----
clang@6.0
```

compilers.yaml

```
compilers:
- compiler:
  modules: []
  operating_system: ubuntu14
  paths:
    cc: /usr/bin/gcc/4.9.3/gcc
    cxx: /usr/bin/gcc/4.9.3/g++
    f77: /usr/bin/gcc/4.9.3/gfortran
    fc: /usr/bin/gcc/4.9.3/gfortran
    spec: gcc@4.9.3
- compiler:
  modules: []
  operating_system: ubuntu14
  paths:
    cc: /usr/bin/clang/6.0/clang
    cxx: /usr/bin/clang/6.0/clang++
    f77: null
    fc: null
    spec: clang@6.0
- compiler:
  ...
```

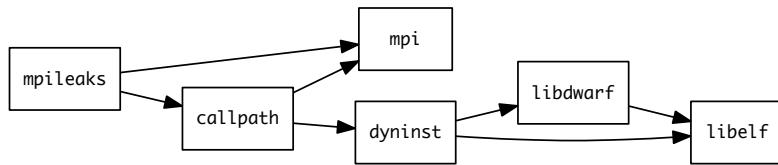
Core Spack Concepts

Most existing tools do not support combinatorial versioning

- Traditional binary package managers
 - RPM, yum, APT, yast, etc.
 - Designed to manage a single stack.
 - Install *one* version of each package in a single prefix (/usr).
 - Seamless upgrades to a *stable, well tested* stack
- Port systems
 - BSD Ports, portage, Macports, Homebrew, Gentoo, etc.
 - Minimal support for builds parameterized by compilers, dependency versions.
- Virtual Machines and Linux Containers (Docker)
 - Containers allow users to build environments for different applications.
 - Does not solve the build problem (someone has to build the image)
 - Performance, security, and upgrade issues prevent widespread HPC deployment.

Spack handles combinatorial software complexity.

Dependency DAG



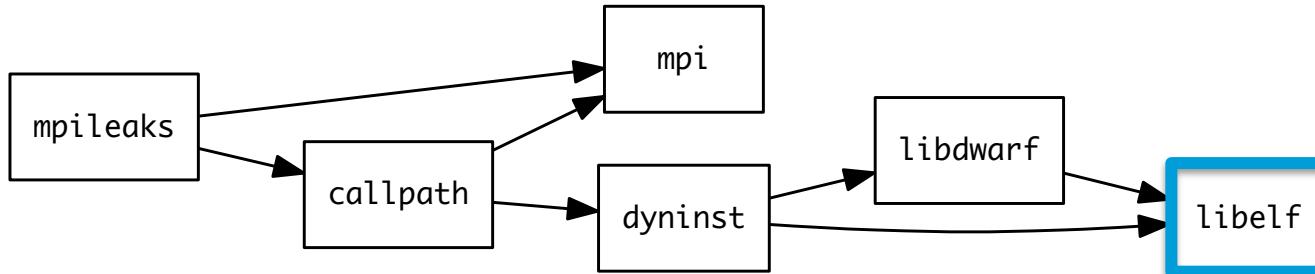
Installation Layout

```
spack/opt/
linux-x86_64/
gcc-4.7.2/
mpileaks-1.1-0f54bf34cadk/
intel-14.1/
hdf5-1.8.15-1kf14aq3nqiz/
bgq/
xl-12.1/
hdf5-1.8.16-fqb3a15abrxw/
...
```

Hash

- Each unique dependency graph is a unique **configuration**.
- Each configuration installed in a unique directory.
 - Configurations of the same package can coexist.
- **Hash** of entire directed acyclic graph (DAG) is appended to each prefix.
- Installed packages automatically find dependencies
 - Spack embeds RPATHs in binaries.
 - No need to use modules or set LD_LIBRARY_PATH
 - Things work *the way you built them*

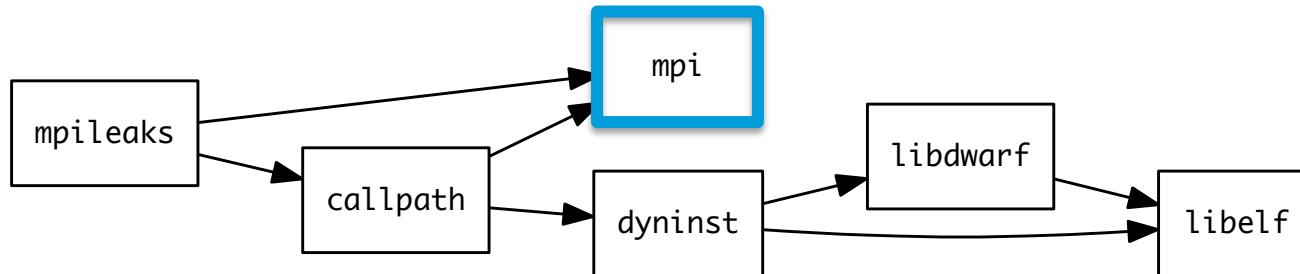
Spack Specs can constrain versions of dependencies



```
$ spack install mpileaks %intel@12.1 ^libelf@0.8.12
```

- Spack ensures *one* configuration of each library per DAG
 - Ensures ABI consistency.
 - User does not need to know DAG structure; only the dependency *names*.
- Spack can ensure that builds use the same compiler, or you can mix
 - Working on ensuring ABI compatibility when compilers are mixed.

Spack handles ABI-incompatible, versioned interfaces like MPI



- *mpi* is a *virtual dependency*
- Install the same package built with two different MPI implementations:

```
$ spack install mpileaks ^mvapich@1.9
```

```
$ spack install mpileaks ^openmpi@1.4:
```

- Let Spack choose MPI implementation, as long as it provides MPI 2 interface:

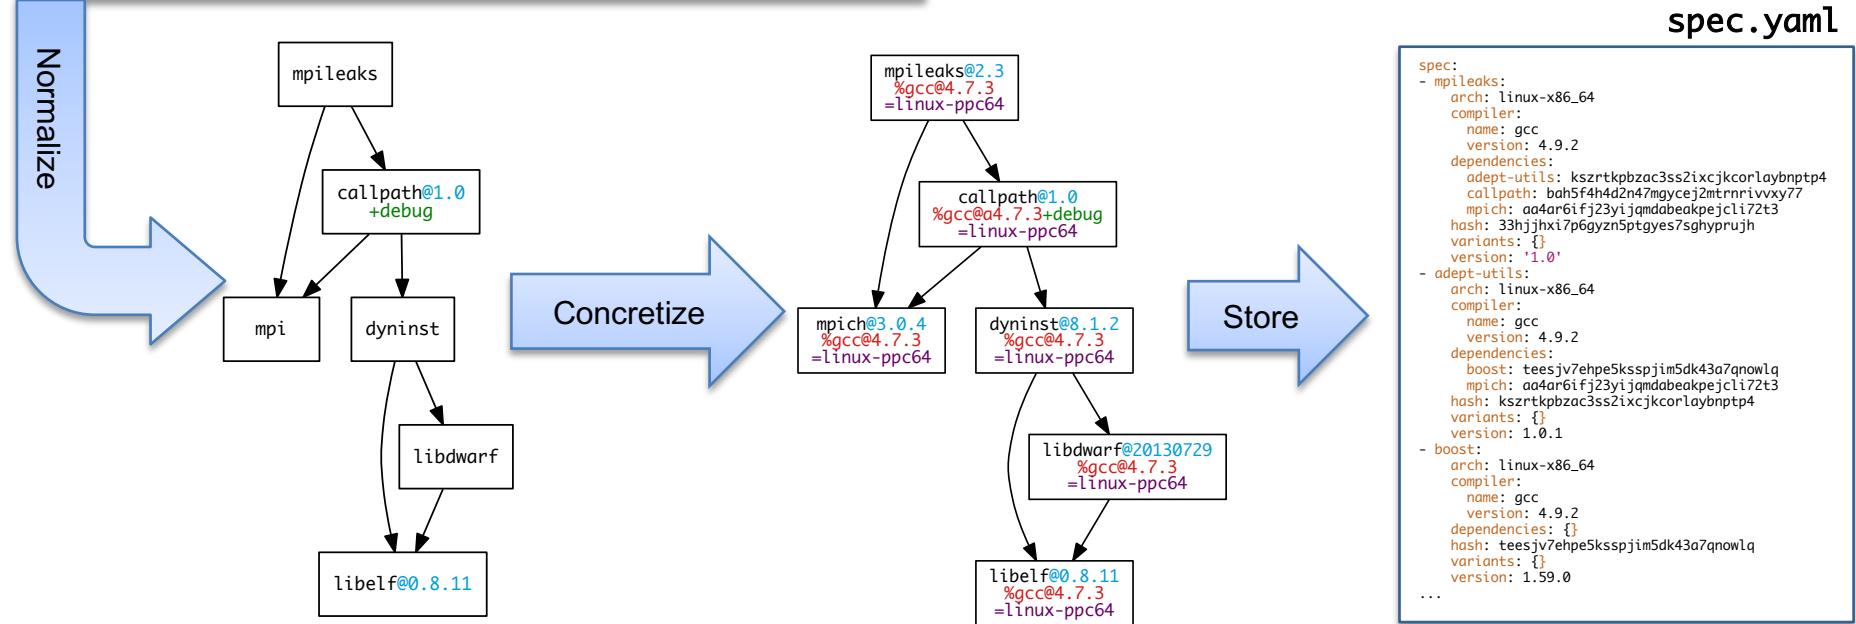
```
$ spack install mpileaks ^mpi@2
```

Concretization fills in missing configuration details when the user is not explicit.

`mpileaks ^callpath@1.0+debug ^libelf@0.8.11`

User input: abstract spec with some constraints

`spec.yaml`



Abstract, normalized spec with some dependencies.

github.com/spack/spack

Use `spack spec` to see the results of concretization

```
$ spack spec mpileaks
Input spec
-----
mpileaks

Normalized
-----
mpileaks
  ^adept-utils
    ^boost@1.42:
      ^mpi
      ^callpath
        ^dyninst
          ^libdwarf
            ^libelf

Concretized
-----
mpileaks@1.0%gcc@5.3.0 arch=darwin-elcapitan-x86_64
  ^adept-utils@1.0.1%gcc@5.3.0 arch=darwin-elcapitan-x86_64
    ^boost@1.61.0%gcc@5.3.0+atomic+chrono+date_time~debug+filesystem~graph
      ~icu_support+iostreams+locale+log+math+mpi+multithreaded+program_options
      ~python+random +regex+serialization+shared+signals+singlethreaded+system
      +test+thread+timer+wave arch=darwin-elcapitan-x86_64
    ^bzzip2@1.0.6%gcc@5.3.0 arch=darwin-elcapitan-x86_64
    ^zlib@1.2.8%gcc@5.3.0 arch=darwin-elcapitan-x86_64
  ^openmpi@2.0.0%gcc@5.3.0~mxm+pmi+psm2~slurm~sqlite3~thread_multiple~tm~verbs+vt arch=darwin-elcapitan-x86_64
  ^hwloc@1.11.3%gcc@5.3.0 arch=darwin-elcapitan-x86_64
    ^libpciaccess@0.13.4%gcc@5.3.0 arch=darwin-elcapitan-x86_64
    ^libtinfo@2.4.6%gcc@5.3.0 arch=darwin-elcapitan-x86_64
    ^m4@1.4.17%gcc@5.3.0+sigsegv arch=darwin-elcapitan-x86_64
    ^libsigsegv@2.10%gcc@5.3.0 arch=darwin-elcapitan-x86_64
^callpath@1.0.2%gcc@5.3.0 arch=darwin-elcapitan-x86_64
  ^dyninst@0.2.0%gcc@5.3.0~stat_dysect arch=darwin-elcapitan-x86_64
  ^libdwarf@20160507%gcc@5.3.0 arch=darwin-elcapitan-x86_64
    ^libelf@0.8.13%gcc@5.3.0 arch=darwin-elcapitan-x86_64
```

Extensions and Python Support

- Spack installs packages in own prefix
- Some packages need to be installed within directory structure of other packages
 - i.e., Python modules installed in \$prefix/lib/python-<version>/site-packages
 - Spack supports this via extensions

```
class PyNumpy(Package):
    """NumPy is the fundamental package for scientific computing with Python."""

    homepage = "https://numpy.org"
    url      = "https://pypi.python.org/packages/source/n/numpy/numpy-1.9.1.tar.gz"
    version('1.9.1', '78842b73560ec378142665e712ae4ad9')

    extends('python')

    def install(self, spec, prefix):
        setup_py("install", "--prefix={0}".format(prefix))
```

Spack extensions

- Examples of extension packages:
 - python libraries are a good example
 - R, Lua, perl
 - Need to maintain combinatorial versioning

```
spack/opt/  
linux-rhel6-x86_64/  
gcc-4.7.2/  
python-2.7.12-6y6vvaw/  
python/  
py-numpy-1.10.4-oaxix36/  
py-numpy/  
...
```

```
$ spack activate py-numpy @1.10.4
```

- Symbolic link to Spack install location
- Automatically activate for correct version of dependency
 - Provenance information from DAG
 - Activate all dependencies that are extensions as well

```
spack/opt/  
linux-rhel6-x86_64/  
gcc-4.7.2/  
python-2.7.12-/  
python/  
py-numpy  
py-numpy-1.10.4-oaxix36/  
py-numpy/
```

Extensions must be activated into extender

```
$ spack extensions python
==> python@2.7.12% gcc@4.9.3 ~tk~ucs4 arch=linux-rhel7-x86_64-i25k4oi
==> 118 extensions:
...
py-autopackage     py-docutils      py-mako          py-numpy        py-py2neo
...
==> 3 installed:
-- linux-rhel7-x86_64 / gcc@4.9.3 -----
py-nose@1.3.7    py-numpy@1.11.0   py-setuptools@20.7.0

==> None activated.

$ spack load python
$ python
Python 2.7.12 (default, Aug 26 2016, 15:12:42)
[GCC 4.9.3] on linux2
Type "help", "copyright", "credits" or "license" for more information.
>>> import numpy
Traceback (most recent call last): File "<stdin>", line 1, in <module>
ImportError: No module named numpy
>>>
```

```
$ spack activate py-numpy
==> Activated extension py-numpy@1.11.0% gcc@4.9.3 +blas+lapack
arch=linux-rhel7-x86_64-77im5ny for python@2.7.12 ~tk~ucs4 % gcc@4.9.3
$ spack extensions python
==> python@2.7.12% gcc@4.9.3 ~tk~ucs4 arch=linux-rhel7-x86_64-i25k4oi
...
==> 1 currently activated:
-- linux-rhel7-x86_64 / gcc@4.9.3 ----

$ python
Python 2.7.12 (default, Aug 26 2016, 15:12:42)
[GCC 4.9.3] on linux2
Type "help", "copyright", "credits" or "license" for more information.
>>> import numpy
>>> x = numpy.zeros(10)
>>>
```

- Python unaware of numpy installation

- activate symlinks entire numpy prefix into Python installation
- Can alternatively load extension

```
$ spack load python
$ spack load py-numpy
```

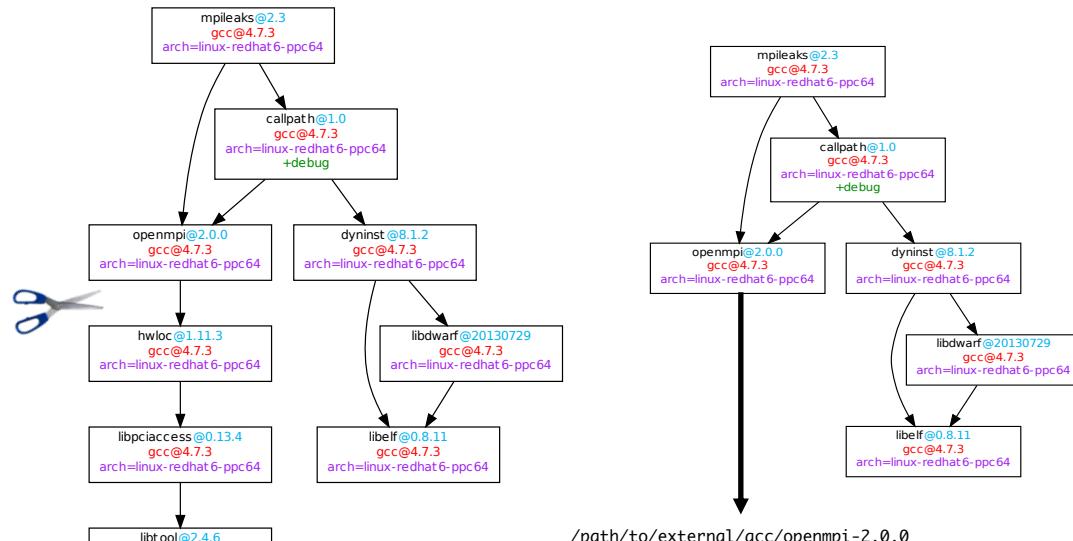
Building against externally installed software

```
mpileaks ^callpath@1.0+debug  
^openmpi ^libelf@0.8.11
```

packages.yaml

```
packages:  
  mpi:  
    buildable: False  
  openmpi:  
    buildable: False  
    paths:  
      openmpi@2.0.0 %gcc@4.7.3 arch=linux-rhel6-ppc64:  
        /path/to/external/gcc/openmpi-2.0.0  
      openmpi@1.10.3 %gcc@4.7.3 arch=linux-rhel6-ppc64:  
        /path/to/external/gcc/openmpi-1.10.3  
      openmpi@2.0.0 %intel@16.0.0 arch=linux-rhel6-ppc64:  
        /path/to/external/intel/openmpi-2.0.0  
      openmpi@1.10.3 %intel@16.0.0 arch=linux-rhel6-ppc64:  
        /path/to/external/intel/openmpi-1.10.3  
...  
...
```

A user registers external packages with Spack.



If a node in the DAG matches an registered external package, Spack prunes the DAG at that node and replaces the node with a reference to the external package.

Spack package repositories

- Some packages can not be released publicly
- Some users have use cases that require ~~bizarre~~ custom builds
- Packaging issues should not prevent users from updating Spack
 - Solution: separate repositories
 - A repository is simply a directory of package files
- Spack supports external repositories that can be layered on top of the built-in packages
- Custom packages can depend on built-in packages (or packages in other repositories)

Proprietary packages

Pathological build cases

“Standard” packages

my_repo

var/spack/repos/builtin

Fetching source code for spack

spack install mpileaks

- User Cache
 - Users create a “mirror” of tar archives for packages
 - Remove internet dependence
- Spack Cache
 - Spack automatically caches tar archives for previously installed software
- The Internet
 - Spack packages can find source files online

Load package from repository

Concretize

Recursively install dependencies

Fetch package source

Build software

User Cache

Spack Cache

The Internet

Adding custom compiler flags

```
$ spack install hdf5 cflags='-O3 -g -fast -fpack-struct'
```

- This installs HDF5 with the specified flags
 - Flags are injected via Spack's compiler wrappers (discussed later).
- Flags are propagated to dependencies automatically
 - Flags are included in the **DAG hash**
 - Each build with different flags is considered a **different version**
- This provides an easy harness for doing parameter studies for tuning codes
 - Previously working with large codes was very tedious.
- Supports cflags, cxxflags, fflags, cppflags, ldflags, and ldlibs
 - Added from CLI or config file

Hands-on Time: Configuration

Follow script at <http://spack.rtfd.io>
Under “Tutorial: Spack 101”

Making your own Spack Packages

Creating your own Spack Packages

- Package is a recipe for building
- Each package is a Python class
 - Download information
 - Versions/Checksums
 - Build options
 - Dependencies
 - Build instructions
- Package collections are repos
 - Spack has a “builtin” repo in `var/spack/repos/builtin`

`$REPO/packages/zlib/package.py`

```
from spack import *

class Zlib(Package):
    """A free, general-purpose, legally unencumbered lossless
       data-compression library."""

    homepage = "http://zlib.net"
    url      = "http://zlib.net/zlib-1.2.8.tar.gz"

    version('1.2.8', '44d667c142d7cda120332623eab69f40')

    depends_on('cmake', type='build')

    def install(self, spec, prefix):
        configure('--prefix={0}'.format(prefix))

        make()
        make('install')
```

Spack packages are *templates* for builds

- Each package has one class
 - zlib for Intel compiler and zlib for GCC compiler are built with the same recipe.
- Can add conditional logic using spec syntax
 - Think of package as *translating* a concrete DAG to build instructions.
 - Dependencies are already built
 - No searching or testing; just do what the DAG says
- Compiler wrappers handle many details automatically.
 - Spack feeds compiler wrappers to (cc, c++, f90, ...) to autoconf, cmake, gmake, ...
 - Wrappers select compilers, dependencies, and options under the hood.

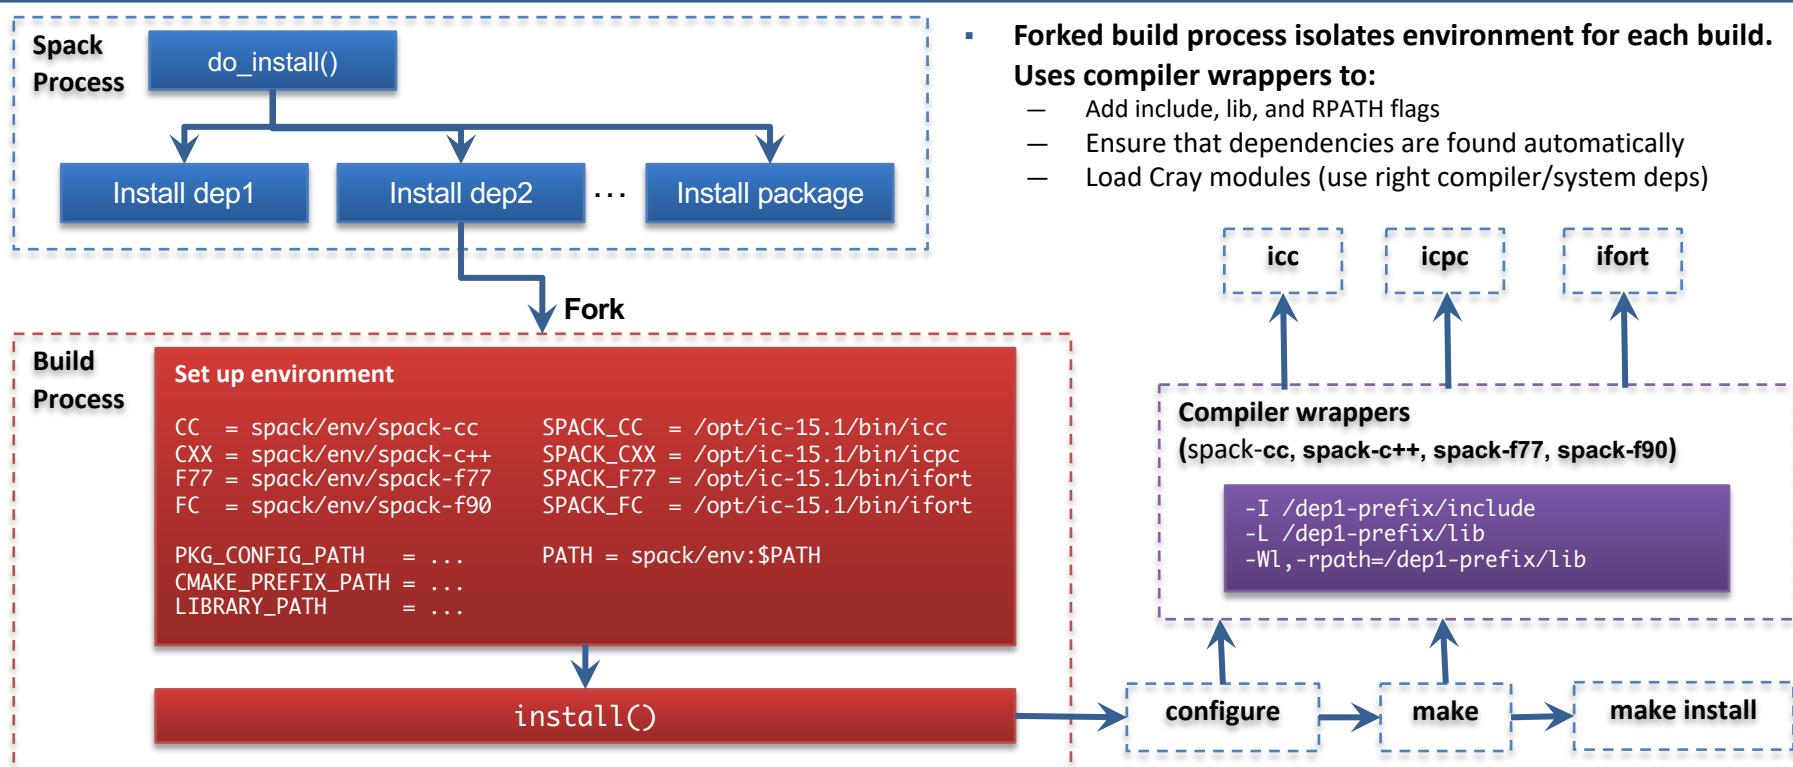
package.py

```
def install(self, spec, prefix):
    config_opts=['--prefix=' + prefix]

    if '~shared' in self.spec:
        config_opts.append('--disable-shared')
    else:
        config_opts.append('--enable-shared')

    configure(config_opts)
    make()
    make('install')
```

Spack builds each package in its own compilation environment



Writing Packages - Versions and URLs

\$REPO/packages/mvapich/package.py

```
class Mvapich2(Package):
    homepage = "http://mvapich.cse.ohio-state.edu/"
    url = "http://mvapich.cse.ohio-state.edu/download/mvapich/mv2/mvapich2-2.2rc2.tar.gz"

    version('2.2rc2', 'f9082ffc3b853ad1b908cf7f169aa878')
    version('2.2b',   '5651e8b7a72d7c77ca68da48f3a5d108')
    version('2.2a',   'b8ceb4fc5f5a97add9b3ff1b9cbe39d2')
    version('2.1',    '0095ceecb19bbb7fb262131cb9c2cdd6')
```

- Package downloads are hashed with MD5 by default
 - Also supports SHA-1, SHA-256, SHA-512
 - We'll be switching to SHA-256 or higher soon.
- Download URLs can be automatically extrapolated from URL.
 - Extra options can be provided if Spack can't extrapolate URLs
- Options can also be provided to fetch from VCS repositories

Writing Packages – Variants and Dependencies

\$REPO/packages/petsc/package.py

```
class Petsc(Package):
    variant('mpi',      default=True,  description='Activates MPI support')
    variant('complex',  default=False, description='Build with complex numbers')
    variant('hdf5',     default=True,  description='Activates support for HDF5 (only parallel)')

    depends_on('blas')
    depends_on('python@2.6:2.7')

    depends_on('mpi', when='+mpi')
```

- Variants are named, have default values and help text
- Other packages can be dependencies
 - **when** clause provides conditional dependencies
 - Can depend on specific versions or other variants

Writing Packages – Build Recipes

- Functions wrap common ops
 - cmake, configure, patch, make, ...
 - **Executable** and **which** for new wrappers.
- Commands executed in clean environment
- Full Python functionality
 - Patch up source code
 - Make files and directories
 - Calculate flags
 - ...

\$REPO/packages/dyninst/package.py

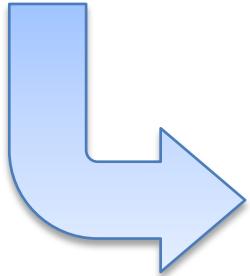
```
def install(self, spec, prefix):
    with working_dir("build", create=True):
        cmake("../", *std_cmake_args)
        make()
        make("install")

@when('@:8.1')
def install(self, spec, prefix):
    configure("--prefix=" + prefix)
    make()
    make("install")
```

Create new packages with spack create

```
$ spack create http://zlib.net/zlib-1.2.8.tar.gz
```

\$REPO/packages/zlib/package.py



```
class Zlib(Package):
    # FIXME: Add a proper url for your package's homepage here.
    homepage = "http://www.example.com"
    url      = "http://zlib.net/zlib-1.2.8.tar.gz"
    version('1.2.8', '44d667c142d7cda120332623eab69f40')

    def install(self, spec, prefix):
        # FIXME: Modify the cmake line to suit your build system here.
```

- `spack create <url>` will create a skeleton for a package
 - Spack reasons about URL, hash, version, build recipe.
 - Generates boilerplate for Cmake, Makefile, autotools, Python, R, Waf, Perl
 - Not intended to completely write the package, but gets you 80% of the way there.
- `spack edit <package>` for subsequent changes

Hands-on Time: Environments

Follow script at <http://spack.rtfd.io>
Under “Tutorial: Spack 101”

Hands-on Time: Module Files

Follow script at <http://spack.rtfd.io>
Under “Tutorial: Spack 101”

Hands-on Time: Creating Packages

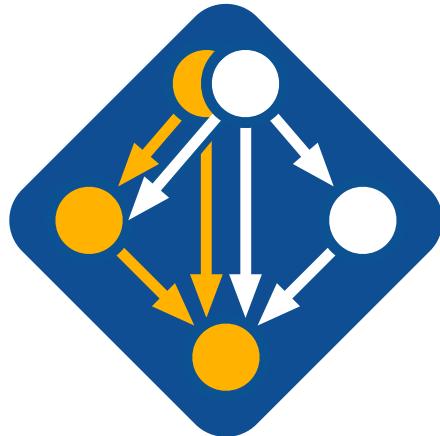
Follow script at <http://spack.rtfd.io>
Under “Tutorial: Spack 101”

Hands-on Time: Advanced Packaging

Follow script at <http://spack.rtfd.io>
Under “Tutorial: Spack 101”

Join the Spack community!

- Contributing packages to Spack is simple
 - Make packages on your own system
 - Send a pull request to Spack to let others use them
 - GitHub guide to pull requests:
<https://help.github.com/articles/creating-a-pull-request/>
 - See contributor guide in the Spack repository
 - Spack is licensed under **MIT** or **Apache 2.0**
- We want more than just packages
 - New features
 - New documentation
 - Any contribution can help the community!
- Spack has a helpful online community
 - Typically happy to respond to GitHub issues
 - Active mailing list on Google Groups



github.com/spack/spack



@spackpm

We hope to make distributing & using HPC software easy!



**Lawrence Livermore
National Laboratory**